

Introduction to Real-Time Systems

Prof. Dr. Giovanni Gracioli

`giovani@lisha.ufsc.br`

`http://www.lisha.ufsc.br/Giovani`

Objectives

- Introduction to Real-Time Systems
 - Motivation for studying
 - Formal definition
 - Soft x hard real-time systems
 - Predictability

What you need to know to follow

- Experience with operating systems
 - Scheduling
 - Memory management
 - Resource management


Books

- J. Liu: Real-Time Systems. Prentice Hall, 2000.
 - Best as an intro book, less good as a reference
- Farines, Silva, Oliveira. Sistemas de Tempo Real. DAS/UFSC. Julho 2000
 - Good introduction to single-core real-time systems
- G.C. Buttazzo: Hard Real-Time Computing Systems: Predictable Scheduling Algorithms and applications. Springer, 2004
 - Good reference on real-time scheduling
- Hermann Kopetz. Real-Time Systems: design principles for Distributed Embedded Applications

What you will learn

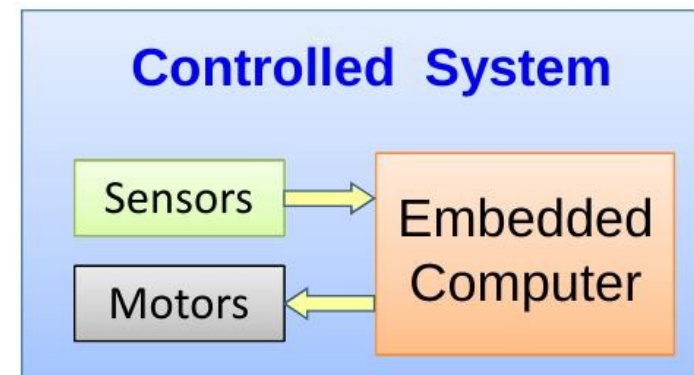
- Definition of real-time systems
- Motivation for studying real-time systems
- Soft x hard real-time systems
- The most important concept of a real-time system
 - Predictability

Let's get started



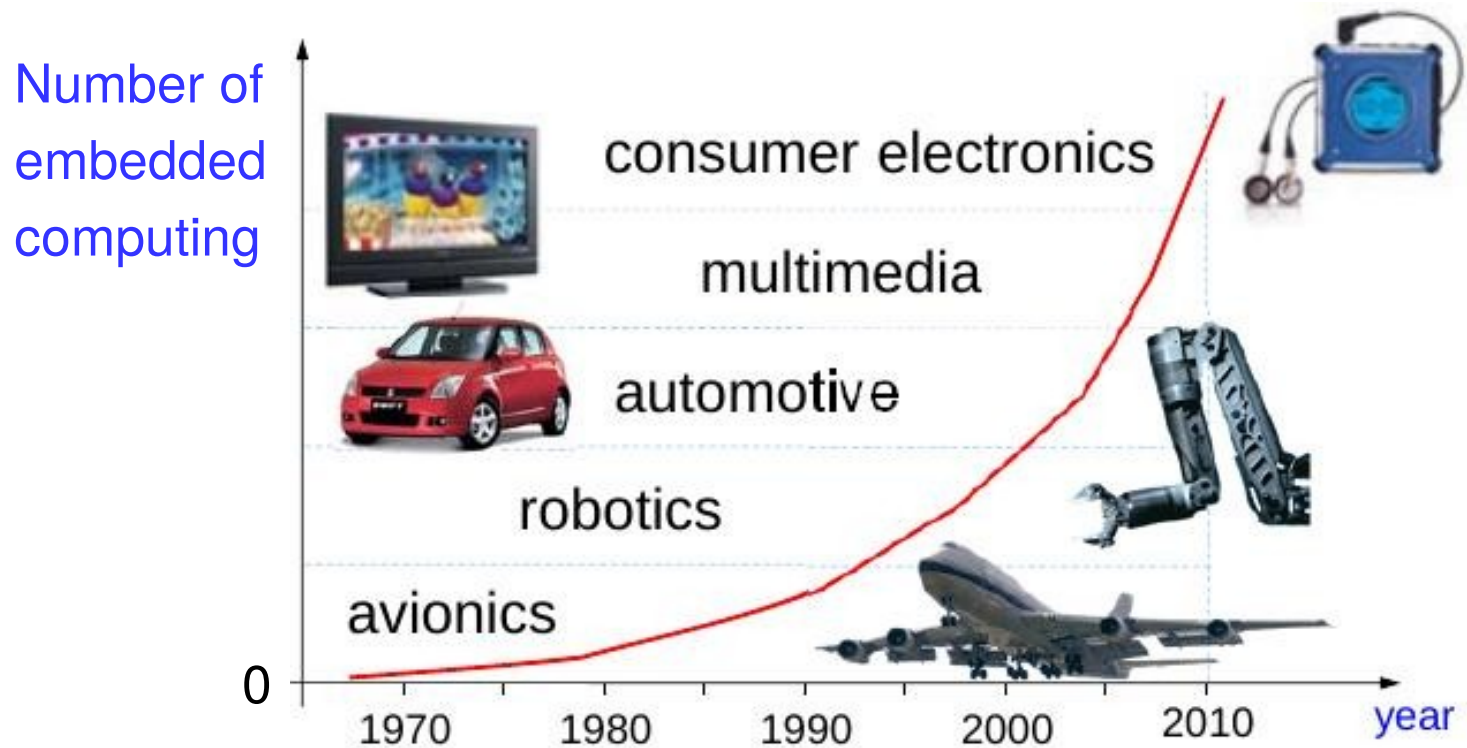
Formal Definition

- The correctness of the system depends not only on the logical result of the computation but also on the time at which the results are produced
 - A correct value at the wrong time is a fault
- They are typically embedded in a larger system to control its functions:
 - Real-Time Embedded Systems



Evolution of Embedded Systems

- Embedded computing systems have grown exponentially in several application domains:




Computers Everywhere

- Today, 98% of all processors in the planet are embedded in other objects



Typical Applications

- avionics
- automotive
- robotics
- industrial automation
- telecommunications
- multimedia systems
- consumer electronics



From Hardware to Software

- We are experiencing a **dematerialization** process in which many functions are converted into software

- Examples

Money

Documents

Books

Music

Pictures


Movies

Tickets

Education



Steer by Wire



Why?

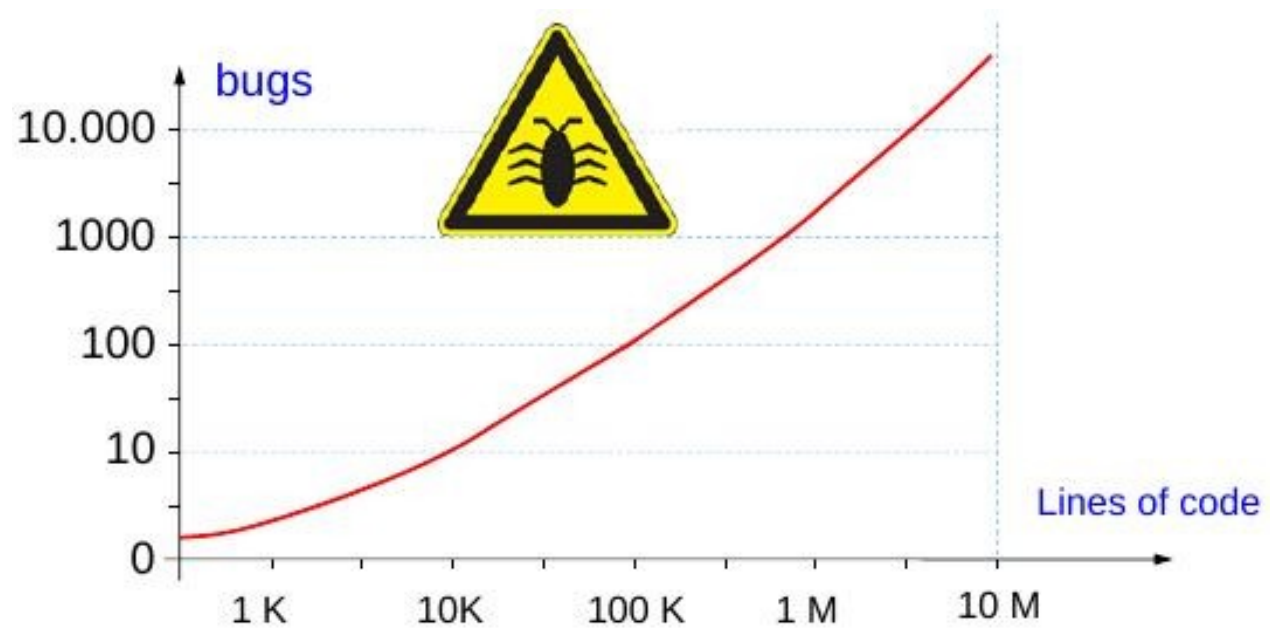
- There are many advantages
 - Software is more flexible than hardware
 - It can be quickly changed/adapted/updated
 - It can be upgraded remotely
 - It can evolve into intelligent control algorithms
 - It has no mass, so it can “travel” at the speed of light

Increasing complexity

- The price to be paid is a higher software complexity
- Related problems
 - Difficult design
 - Less predictability
 - Less reliability
- Novel solutions for
 - Component-based software design
 - Analysis for guaranteeing predictability and safety
 - Testing


Complexity and bugs

- Software bugs increase with complexity



Software reliability

- **Reliability** does not only depend on the correctness of single instructions, but also on **when** they are executed



A correct action executed too late can be **useless** or even **dangerous**

Real-Time Systems

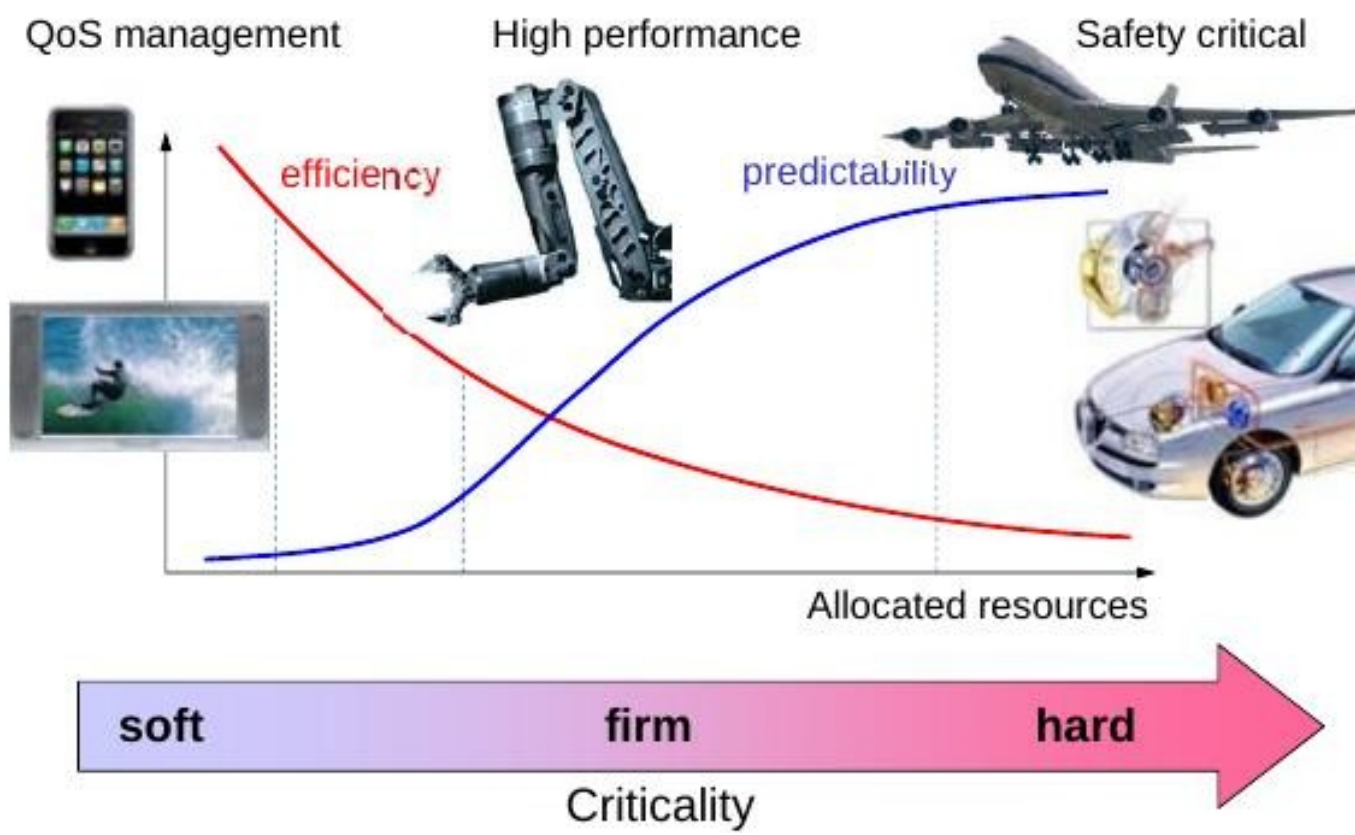
Computing systems that must guarantee **bounded** and **predictable** response times are called **real-time systems**

- Activities are associated with timing constraints (**deadlines**)
- Predictability of response times must be guaranteed
 - for each critical activity
 - for all possible combination of events

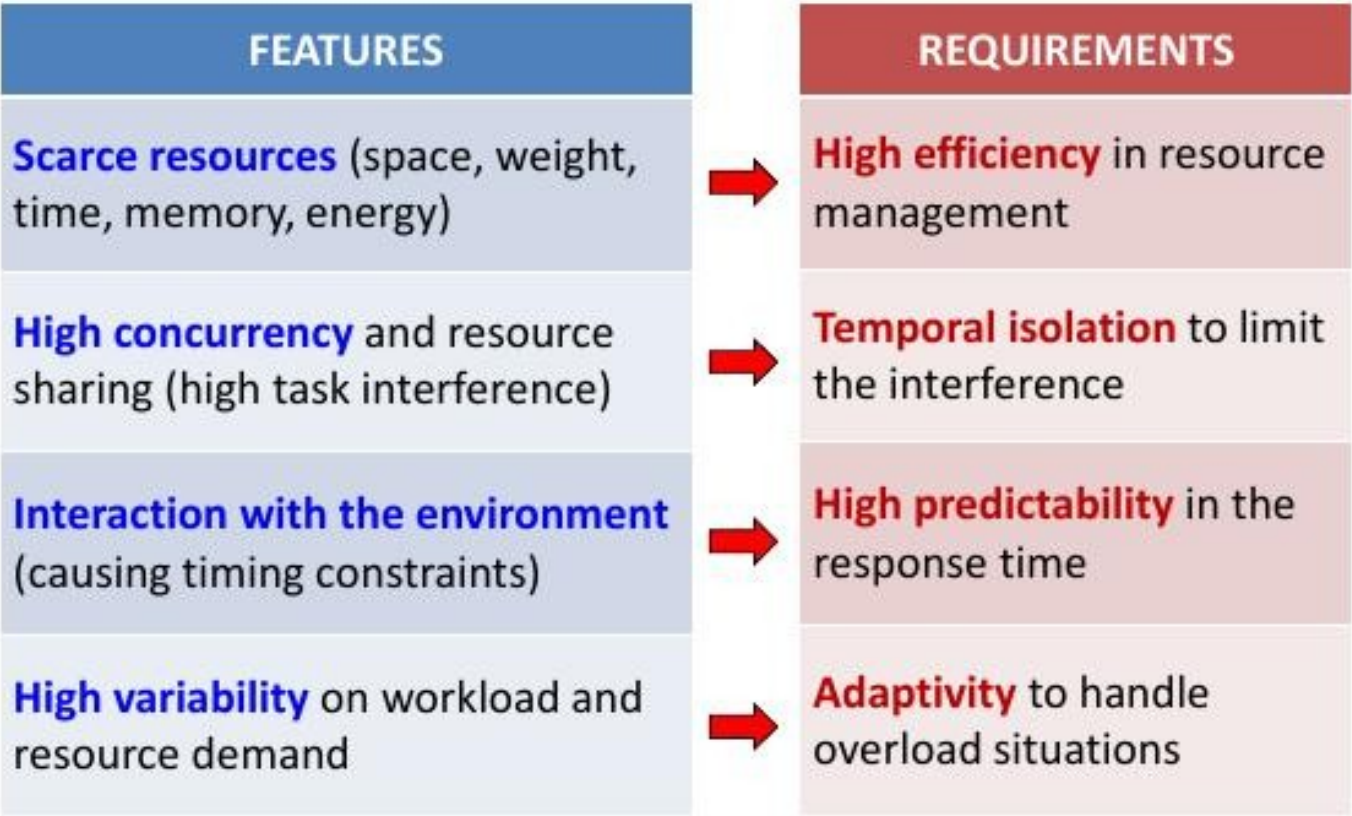
Soft x Hard Real-Time

- **Soft Real-Time (SRT):** missing deadlines is undesirable, but will not lead to catastrophic consequences
 - Related to the concept of “Quality of Service”
 - Typically interested in average-case response time
 - Ex: reservation systems, media players, phones
- **Hard Real-Time (HRT):** missing deadlines is not an option
 - Interested in worst-case response times
 - Ex: airplanes, nuclear plants, military systems

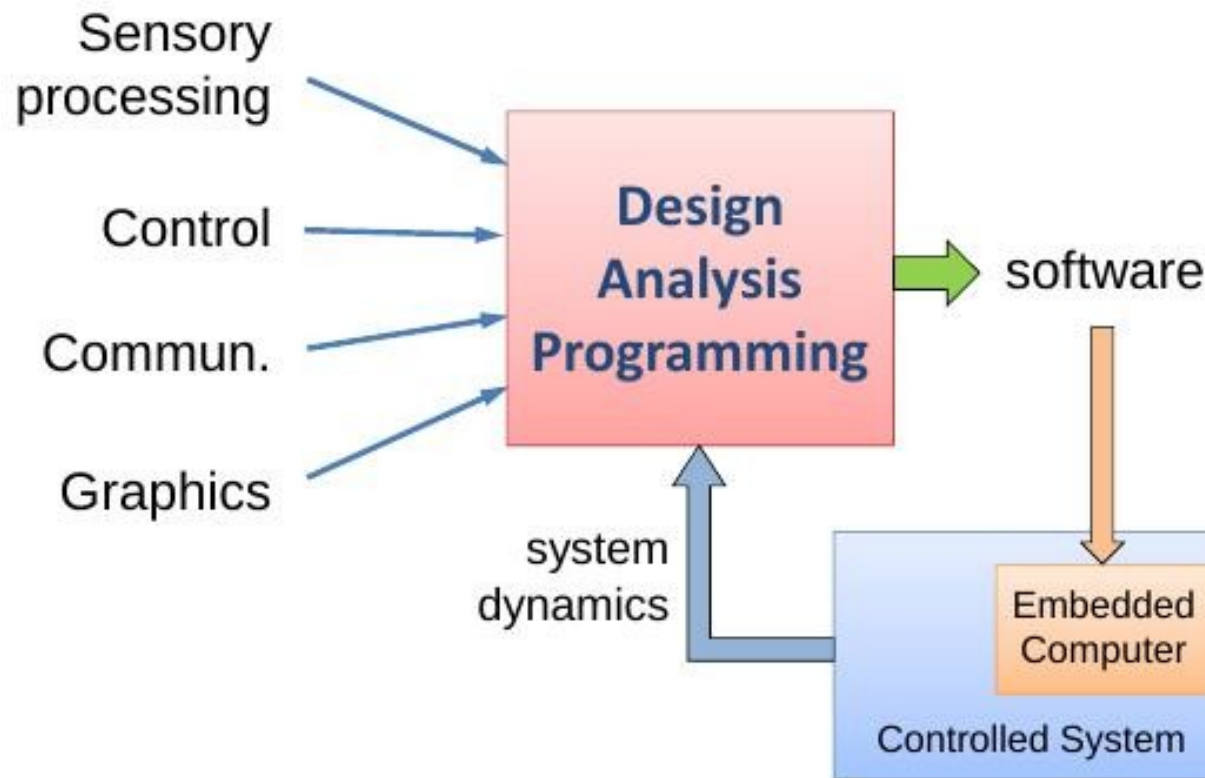
Real-Time Spectrum



What's special in Embedded Systems?



Our focus: predictable software



Control and implementation

- Often, control and implementation are done by different people that do not talk to each other



- Control guys typically assume a computer with infinite resources and computational power. In some cases, computation is modeled by a fixed delay



Control and implementation

- In reality, a computer:
 - has **limited** resources
 - **finite** computational power (**non null execution times**)
 - executes several **concurrent** activities
 - introduces **variable** delays (often **unpredictable**)

Modeling such factors and taking them into account in the design phase allows a significant improvement in performance and reliability

A control example




A positive angle θ requires a positive control action μ .



A control task

```
task control(float theta0, float k)
{
  float error;
  float u;
  float theta;

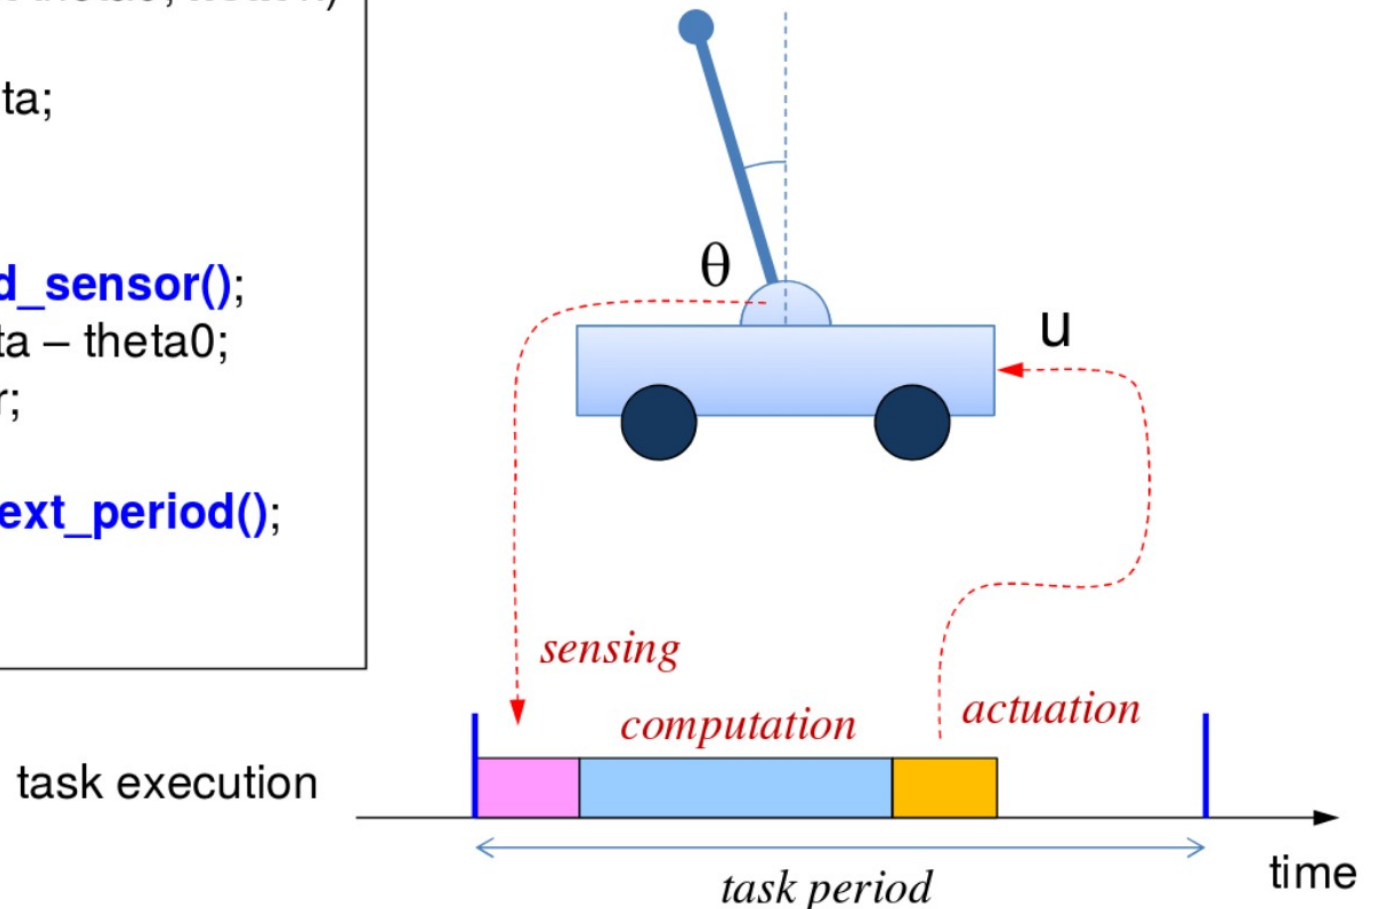
  while (1) {
    theta = read_sensor();
    error = theta - theta0;
    u = k * error;
    output(u);
    wait_for_next_period();
  }
}
```



A control task


```
task control(float theta0, float k)
{
float error, u, theta;

    while (1) {
        theta = read_sensor();
        error = theta - theta0;
        u = k * error;
        output(u);
        wait_for_next_period();
    }
}
```




Tradition control view

Negligible delay and jitter




Effect of computation times

Computation times introduce a non negligible delay:



Actual Situation

Actual situation: variable delay and jitter:



RTOS responsibilities

- A **real-time operating system** (RTOS) is responsible for
 - managing **concurrency**
 - activating periodic tasks at the beginning of each period (**time management**)
 - deciding the execution order of tasks (**scheduling**)
 - solving possible timing conflicts during the access of shared resources (**mutual exclusion**)
 - manage the timely execution of asynchronous events (**interrupt handling**)

Typical objection

It is not worth to invest in RT theory, because computer speed is increasing exponentially, and all timing constraints can eventually be handled

■ Answer

- Given an arbitrary computer speed, we must always guarantee that timing constraints can be met. Testing is NOT sufficient.

Real-Time != Fast

- A real-time system **is not** a fast system
- Speed is always relative to a specific environment
- Running fast is good, but does not guarantee a correct behavior

Speed vs. Predictability

- The objective of a real-time system is to guarantee the timing behavior of each individual task
- The objective of a fast system is to minimize the average response time of a task set. However,

Do not trust the **average** when you have to guarantee individual performance

Sources of non determinism

- **Architecture**
 - cache, pipelining, interrupts, DMA
- **Operating System**
 - scheduling, synchronization, communication
- **Language**
 - lack of explicit support for time
- **Design methodologies**
 - lack of analysis and verification techniques

Traditional (wrong) approach

- Most RT application are designed using empirical techniques
 - assembly programming
 - timing through dedicated timers
 - control through driver programming
 - priority manipulation

Disadvantages

- 1) Tedious programming which heavily depends on programmer's ability
- 2) Difficult code understanding
- 3) Difficult software maintainability
 - a) millions lines of code
 - b) code understanding takes more time than re-writing
 - c) re-writing is very expensive and bug prone
- 4) Difficult to verify timing constraints without explicit support from OS and the language

Implications

- Such a way of programming RT applications is very dangerous
- It may work in most situations, but the risk of a failure is high
- When the system fails is very difficult to understand why
- Conclusion: **low reliability**

Accidents due to SW

- Task overrun during LEM lunar landing
(<http://njnnetwork.com/2009/07/1202-computer-error-almost-aborted-lunar-landing/>)
- First flight of the Space Shuttle (synch)
- Ariane 5 (overflow)
- Airbus 320 (cart task)
- Airbus 320 (holding task)
- Pathfinder (reset for timeout, priority inversion)

Lessons learned

- Tests, although necessary, allow only a partial verification of system's behavior
- Predictability must be improved at the level of the operating system
- The system must be designed to be fault-tolerant and handle overload conditions
- Critical systems must be designed under pessimistic assumptions


and always remember of the Murphy's Law

- If something can go wrong, it will go wrong
- If a system stops working, it will do it at the worst possible time
- Sooner or later, the worst possible combination of circumstances will occur

Proving Murphy's law

Let p_E be the probability for event E to occur in a day

What is the probability for E to occur in n days?



prob. of E not
occurring in 1 day

$$q_E = 1 - p_E$$


prob. of E not
occurring in n days

$$Q_E(n) = (1 - p_E)^n$$

prob. of E
occurring in n day

$$P_E(n) = 1 - Q_E(n)$$

Understanding Murphy's law



If something can go wrong (no matter how small P_e is), it will go wrong (that is, the probability for E to occur in long time intervals tends to 1).

Review

- Formal definition of Real-Time Systems
- Soft x Hard real-time systems
- Predictability

References

- Giorgio Buttazo. Real-time systems course
- Farines, Silva, Oliveira. Sistemas de Tempo Real. DAS/UFSC. Julho 2000
- G.C. Buttazzo: Hard Real-Time Computing Systems: Predictable Scheduling Algorithms and applications. Springer, 2004