

Real-Time Systems: concepts, task models, and uniprocessor scheduling

Prof. Dr. Giovanni Gracioli

`giovani@lisha.ufsc.br`

`http://www.lisha.ufsc.br/Giovani`

Objectives

- Basic concepts of Real-Time Systems
 - Task models
 - the periodic task model
 - Schedulability analysis/schedulability tests
 - Uniprocessor real-time schedulers
 - RM
 - DM
 - EDF
 - LLF


What you need to know to follow

- Experience with operating systems
 - Scheduling
 - Memory management
 - Resource management
- What a RTS is
- Soft/Hard real-time systems

What you will learn

- Periodic real-time model
- Uniprocessors real-time schedulers
 - Cyclic executive
 - RM
 - DM
 - EDF
 - LLF

Let's get started



Task Model

- **Task**: a real-time computation unit
 - Think of it as an execution thread with additional timing parameters
 - The job of a real-time scheduler is to schedule tasks
- Three main task models
 - Aperiodic
 - Periodic
 - Sporadic

Aperiodic Tasks

- Event-triggered computation
- Task is activated by an external event
- Task runs once to respond to the event
- Relative deadline D : available time to respond to the event
- Ex:
 - Event = loss of power
 - Task = drop control rods into nuclear reactor
(this is actually happened in Fukushima)

Periodic Tasks

- Time-triggered computation
- Task is activated **periodically** every **T time units**
- Each periodic instance of the task is called a **job**
- Each job has the same **relative deadline** (usually = to period)
- Ex:
 - most digital controllers

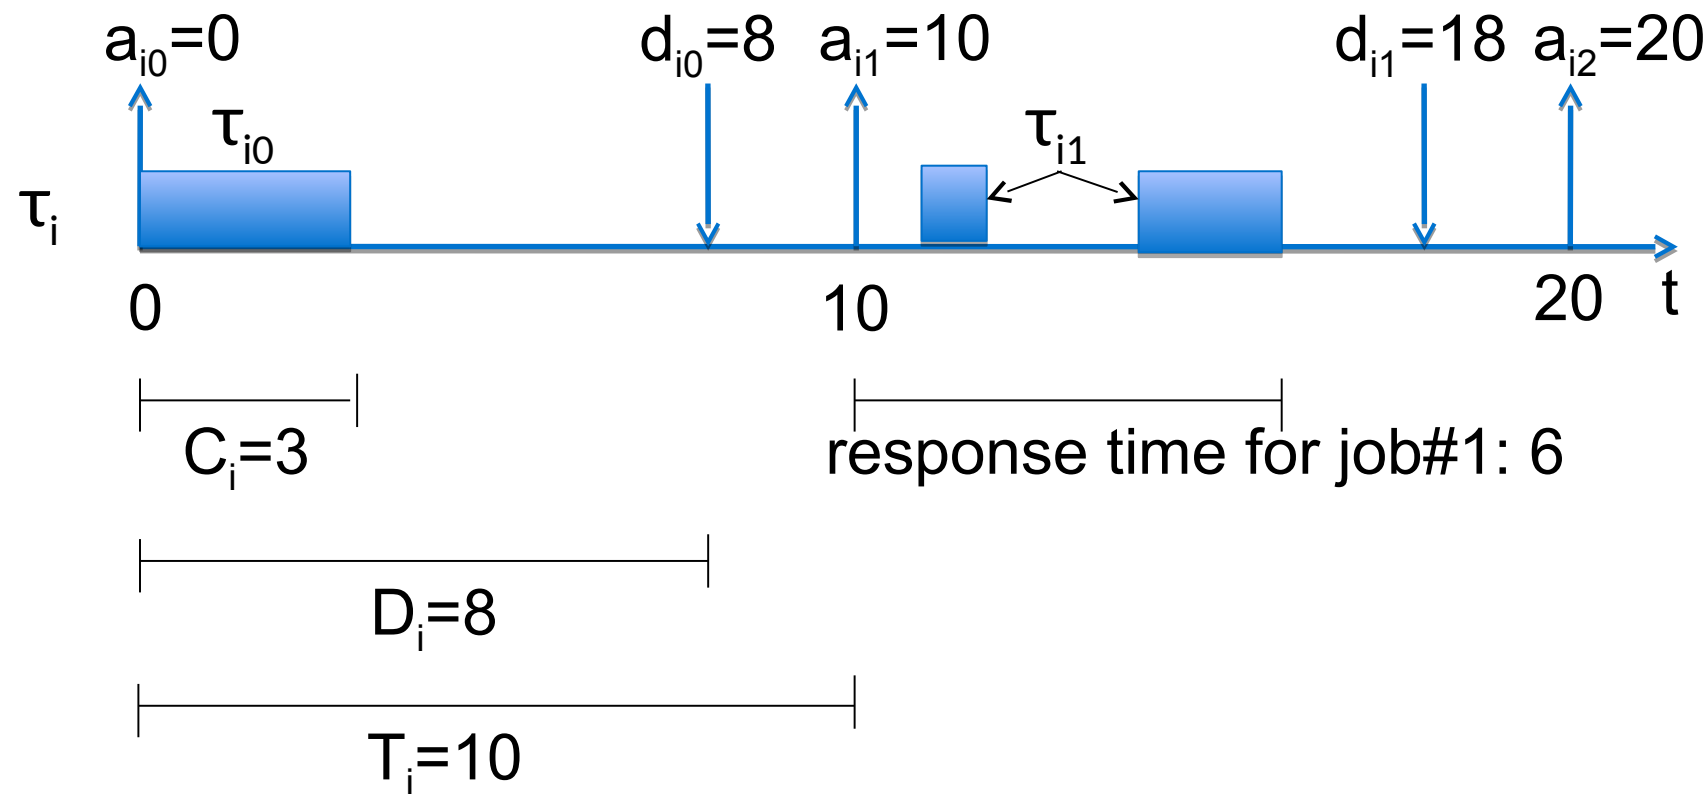
Sporadic Tasks

- Same as periodic task, but the task is activated at most every T time units (minimum inter-arrival time)
- Ex:
 - processing network packets


Periodic Tasks - Main concepts

- Task τ_i (N tasks in the system, τ_1 to τ_N)
 - Execution time C_i (sometimes e_i)
 - Relative deadline D_i
 - Period T_i (sometimes p_i)
- Each job τ_{ij} of τ_i (first job: τ_{i0})
 - Activation time $a_{ij} = a_{ij-1} + T_i$ (usually with $a_{i0} = 0$)
 - Sporadic: $a_{ij} \geq a_{ij-1} + T_i$
 - Absolute deadline $d_{ij} = a_{ij} + D_i$

Periodic Task Model




Periodic Task Model

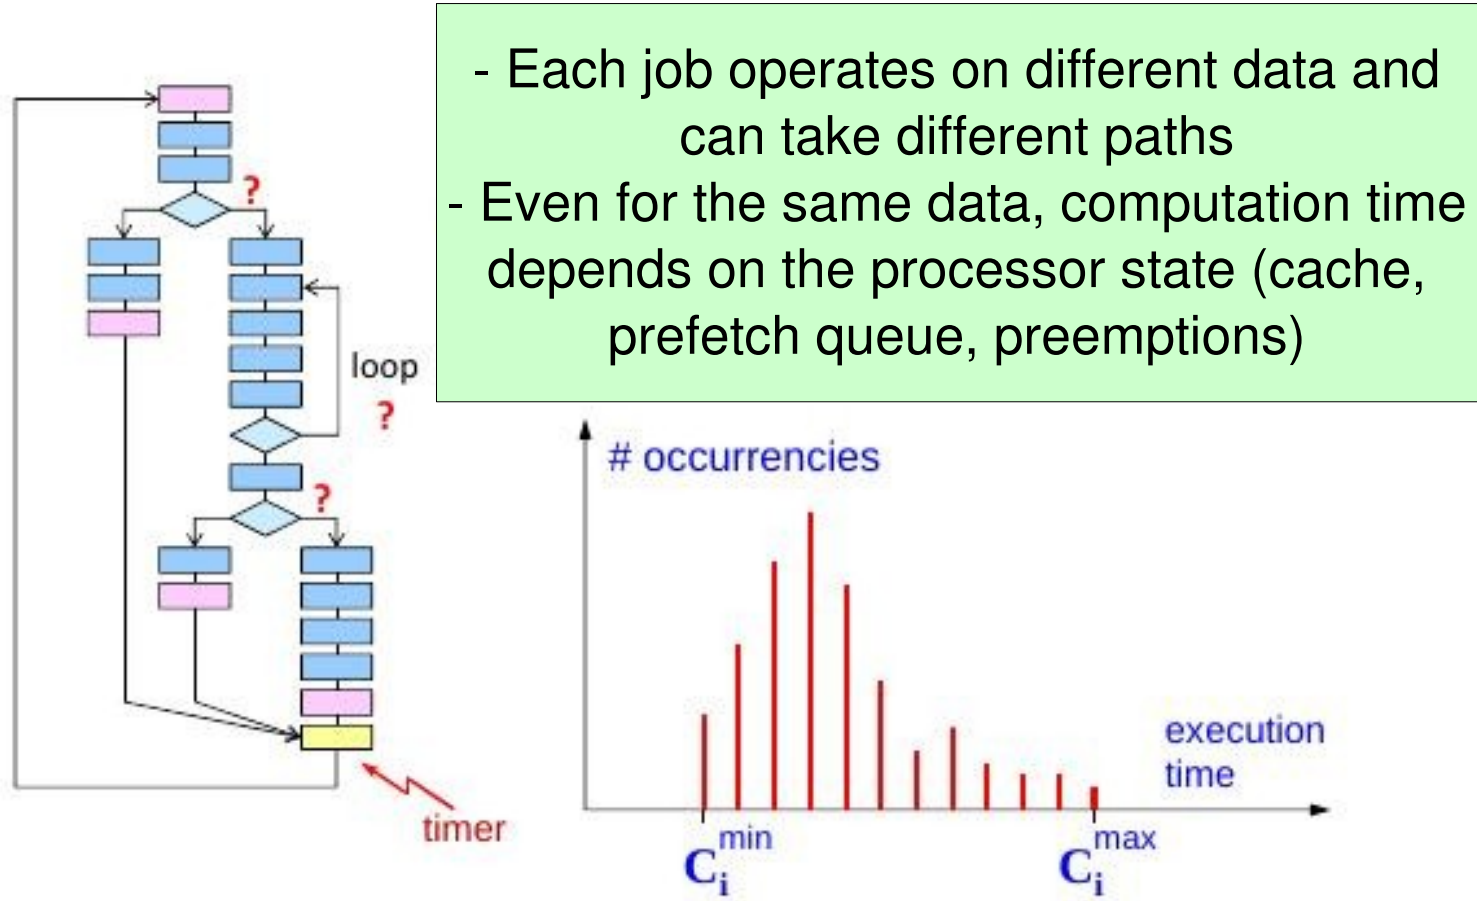


- start time (s_i), finish time (f_i), response time R_i
- A real-time task τ_i is said to be **feasible** if it is guaranteed to complete within its deadline, that is, if $f_i \leq d_i$, or equivalently, if $R_i \leq D_i$


Slack and Lateness



Estimating C_i is not easy




Predictability x Efficiency




Jitter

- It is a measure of the time variation of a periodic event
- Start time jitter
 - Time to handle time interrupt and to insert the job into the running queue



Parameters summary



- Computation time (C_i)
- Period (T_i)
- Relative deadline (D_i)
- Arrival time (a_i)
- Start time (s_i)
- Finishing time (f_i)
- Response time (R_i)
- Slack and Lateness
- Jitter

These parameters are specified by the programmer and are known off-line


These parameters depend on the scheduler and on the actual execution, and are known at run time.

Utilization & Schedulability analysis

- Task Utilization for a periodic/sporadic task:
 $U_i = C_i / T_i$
 - Percentage of processor time required by the task
- System Utilization: $U = U_1 + U_2 + \dots + U_N$
 - Percentage of processor time required by all tasks
- Base uniprocessor scheduling result: task set is clearly not schedulable if: $U > 1$
- For many scheduling algorithms, we can define a utilization bound U_b such that the task set is schedulable if: $U \leq U_b$

Schedulability

- Set of tasks is **schedulable** under a set of constraints, if a schedule exists for that set of tasks & constraints
- **Exact** tests are NP-hard in many situations
- **Sufficient tests**: sufficient conditions for schedule checked. (Hopefully) small probability of not guaranteeing a schedule even though one exists
- **Necessary tests**: checking necessary conditions. Used to show no schedule exists. There may be cases in which no schedule exists & we cannot prove it



Task Constraints


- **Timing constraints**
 - activation, completion, jitter
- **Precedence constraints**
 - they impose an ordering in the execution
- **Resource constraints**
 - they enforce a synchronization in the access of mutually exclusive resources

Real-Time Scheduling

- Assume that we are given a task graph $G=(V,E)$
- **Def.:** A schedule T of G is a mapping

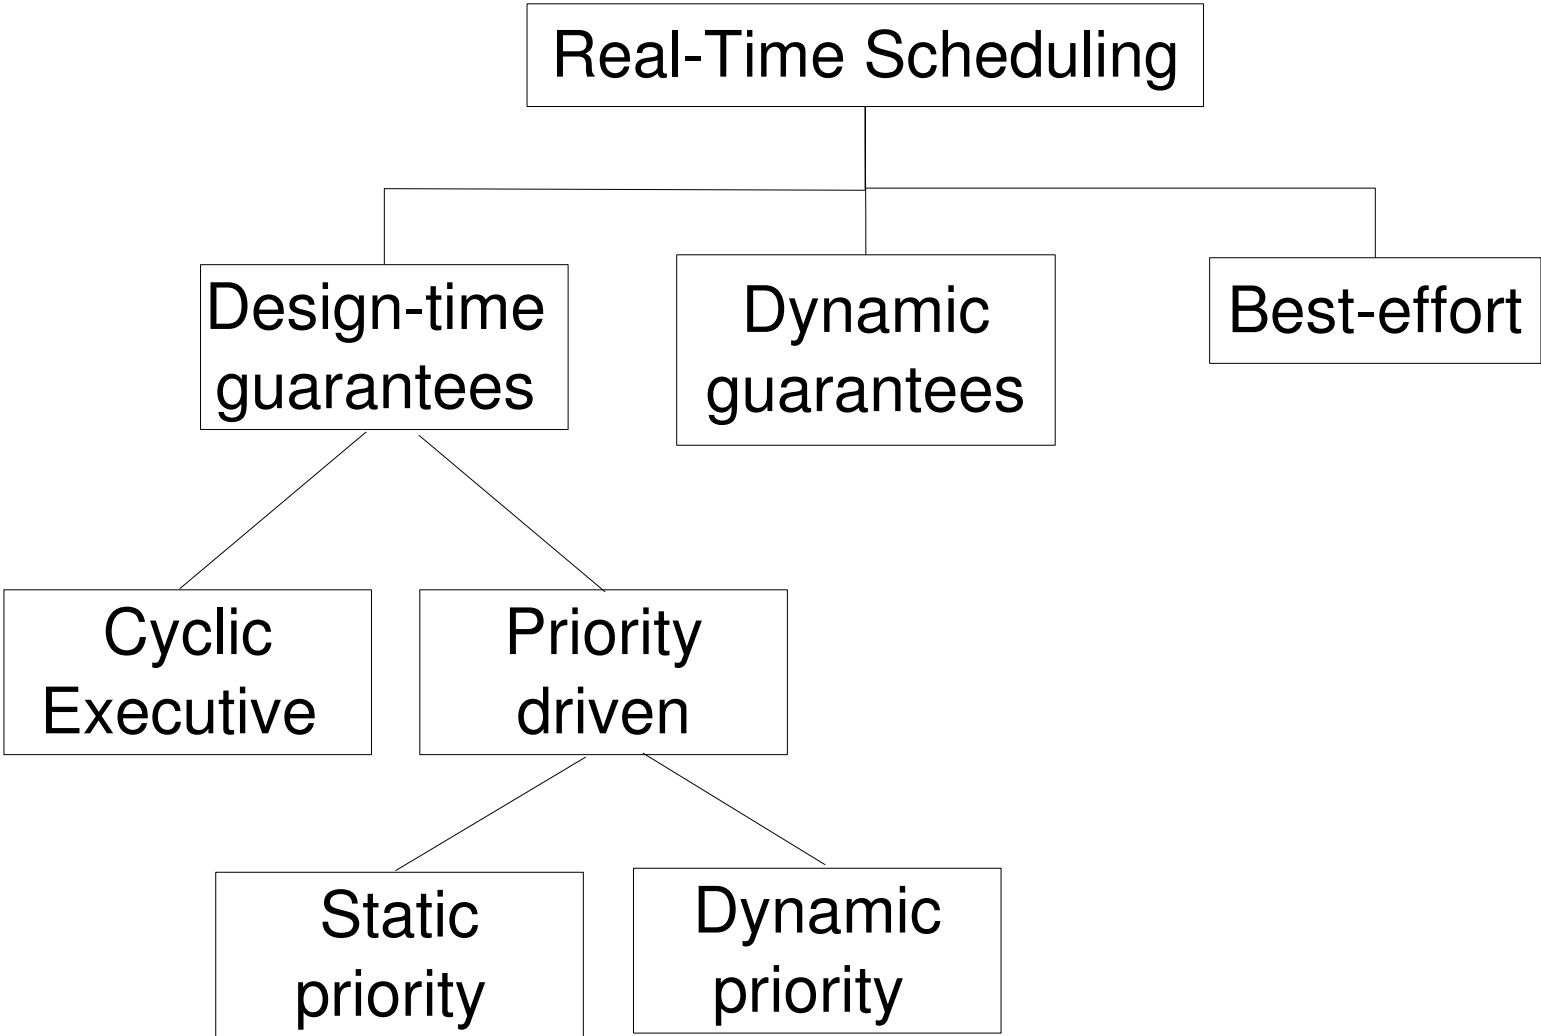
$$V \rightarrow D_t$$

of a set of tasks V to start times from domain D_t



- Typically, schedules have to respect a number of constraints (resource, dependency, timing)
- Scheduling = finding such a mapping

RT Scheduling Classification



The simplest model

- We first look at the simplest possible model, where:
 - All tasks are periodic
 - Single processor
 - Tasks do not share any resource
- Not very realistic, but instructive (we have simple results)!
- We will then (briefly) look at:
 - What happens when you start sharing resources
 - What happens if you schedule a mix of periodic/aperiodic tasks
 - What happens if you use a multiprocessor
 - More complex task models

Uniprocessor real-time scheduling

■ Three main categories

1) Table-driven scheduling

- build a table that dictates when each task executes

2) Fixed-priority scheduling

- Each task is assigned a fixed priority

3) Dynamic-priority scheduling

- Task priority varies at run-time (each job of the task has a different priority)

■ Scheduler is typically preemptive - better system utilization

Cyclic Executive Scheduling

- Also known as cyclic scheduling
- It has been used for 30 years in military systems, navigation, and monitoring systems
- Only works for periodic tasks and off-line analysis
- **Examples**
 - Air traffic control systems
 - Space shuttle
 - Boeing 777
 - Airbus navigation system

Cyclic Executive Scheduling

- The time axis is divided in intervals of equal length (**time slots**)
- Each task is statically allocated in a slot in order to meet the desired request rate
- The execution in each slot is activated by a timer


Cyclic Executive Scheduling

Example:

task	C_i	T_i
A	10 ms	25 ms
B	10 ms	50 ms
C	10 ms	100 ms

$\Delta = \text{GCD}$ (minor cycle)


$T = \text{lcm}$ (major cycle)



Guarantee: $\begin{cases} C_A + C_B \leq \Delta \\ C_A + C_C \leq \Delta \end{cases}$

Cyclic Executive Scheduling

Implementation:



Cyclic Executive Scheduling

Coding:

```
#define MINOR 25 // minor cycle = 25 ms
initialize_timer(MINOR); // interrupt every 25 ms
while (1) {
    sync(); // block until interrupt
    function_A();
    function_B();
    sync(); // block until interrupt
    function_A();
    function_C();
    sync(); // block until interrupt
    function_A();
    function_B();
    sync(); // block until interrupt
    function_A();
}
```

22

Cyclic Executive Scheduling

■ Advantages

- Simple implementation (no RTOS is required)
- Low run-time overhead
- Very low jitter

■ Disadvantages

- Very fragile during overload conditions => domino effect or task abort
- It is difficult to expand the schedule (sensitivity to application changes)
- It is not easy to handle aperiodic activities

Priority (static or dynamic) Scheduling

■ Method

- Assign priorities to each task according to the scheduling policy (static or dynamic)
- Verify the feasibility of the schedule using analytical techniques (schedulability test of the scheduling algorithm)
- Execute tasks on a priority-based RTOS

Rate Monotonic (RM)

- Each task is assigned a fixed priority proportional to its rate [Liu and Layland 73]
 - shorter period = higher priority
- Assumes periodic or sporadic tasks with $D_i = T_i$ on a uniprocessor
- Assumes that the context switch time is zero


Rate Monotonic (RM)

- Sufficient schedulability test
 - Proposed by Liu and Layland in 1973
 - Task set schedulable if
 - $U \leq U_b(N) = N(2^{1/N} - 1)$
 - When all tasks have harmonic periods $U \leq U_b$
- Note that $\lim_{N \rightarrow +\infty} N(2^{1/N} - 1) = \log 2 \approx 0.693$
- What happens if $U_b(N) < U \leq 1$?
 - Nothing can be said according to the analysis
 - Task set might or might not be schedulable

Rate Monotonic (RM)

$$\mu = \sum_{i=1}^n \frac{c_i}{p_i} \leq n(2^{1/n} - 1)$$


$$\lim_{n \rightarrow \infty} (n(2^{1/n} - 1)) = \ln(2)$$



Rate Monotonic (RM)

■ Case 1


- τ_1 ($C_1 = 4$, $T_1 = 8$), high prio, τ_2 ($C_2 = 3$, $T_2 = 12$), low prio
- Utilization: $U = 4/8 + 3/12 = 0.75 < U_b(2) \approx 0.828$
- Schedulability analysis: schedulable



Rate Monotonic (RM)

■ Case 2


- τ_1 ($C_1 = 4$, $T_1 = 8$), high prio, τ_2 ($C_2 = 4$, $T_2 = 12$), low prio
- Utilization: $U = 4/8 + 4/12 \approx 0.833 > U_b(2) \approx 0.828$
- Schedulability analysis: we do not know
- In reality: it is schedulable



Rate Monotonic (RM)

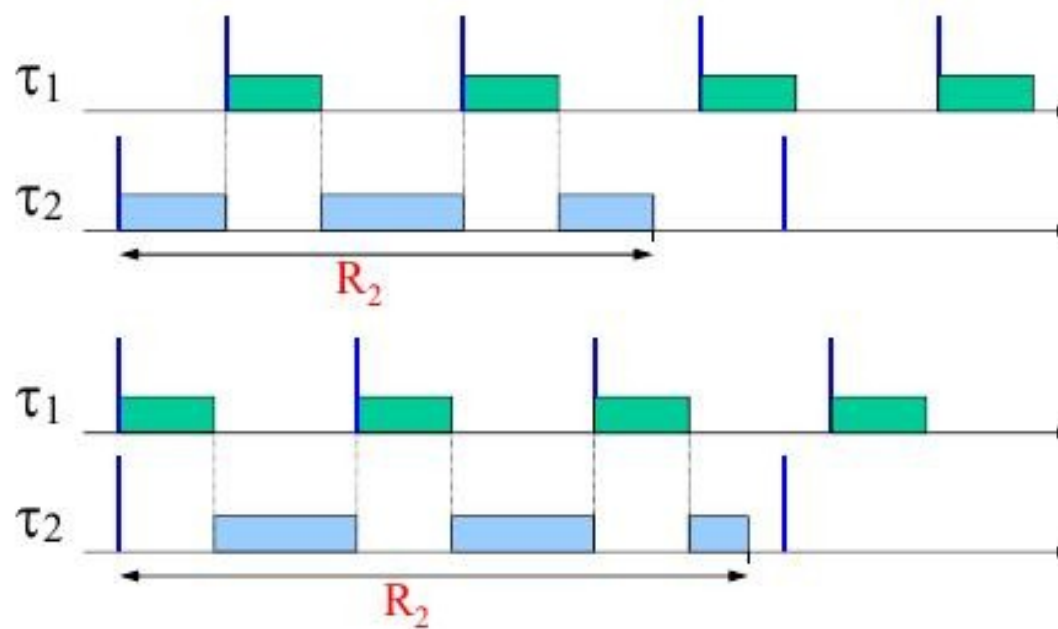
■ Case 3

- τ_1 ($C_1 = 4$, $T_1 = 8$), high prio, τ_2 ($C_2 = 6$, $T_2 = 12$), low prio
- Utilization: $U = 4/8 + 6/12 = 1 > U_b(2) \approx 0.828$
- Schedulability analysis: we do not know
- In reality: it is not schedulable



Critical Instant

- For any task τ_1 the longest response time occurs when it arrives together with all higher priority tasks



Rate Monotonic is optimal

- RM is optimal among all fixed priority algorithms when $D_i = T_i$

if there exists a fixed priority assignment which leads to a feasible schedule, then the RM schedule is feasible



if a task set is not schedulable by RM, then it cannot be schedule by any fixed priority assignment


Earliest-Deadline First (EDF)

- Dynamic-Priority Scheduling Algorithm
- Task priority is inversely proportional to its current absolute deadline
 - Earlier deadline = higher priority
 - Each job of a task has a different deadline, hence a different priority
- Assumes periodic or sporadic tasks with $D_i = T_i$ on a uniprocessor, then...
- Schedulability analysis: the task set is schedulable if: $U \leq U_b = 1$
 - Since task sets with $U > 1$ can not be scheduled by any algorithm, EDF is optimal

EDF example


■ Case 3

- τ_1 ($C_1 = 4$, $T_1 = 8$), high prio, τ_2 ($C_2 = 6$, $T_2 = 12$), low prio
- Utilization: $U = 4/8 + 6/12 = 1$
- Schedulability analysis: schedulable




Another example

- τ_1 ($C_1 = 1, T_1 = 4$), τ_2 ($C_2 = 2, T_2 = 6$), τ_3 ($C_3 = 3, T_3 = 8$)
- Utilization: $U = 1/4 + 2/6 + 3/8 = 23/24$
- RM: don't know (utilization bound), in reality: not schedulable



Another example


- τ_1 ($C_1 = 1, T_1 = 4$), τ_2 ($C_2 = 2, T_2 = 6$), τ_3 ($C_3 = 3, T_3 = 8$)
- Utilization: $U = 1/4 + 2/6 + 3/8 = 23/24$
- EDF: schedulable



EDF is optimal

- EDF is optimal among all algorithms

if there exists a feasible schedule for a task set, then EDF will generate a feasible schedule



if a task set is not schedulable by EDF, then it cannot be scheduled by any algorithm

EDF x RM

- In practice, industrial systems heavily favor RM over EDF. Why?
- For most task sets, RM has better utilization bound than $\log 2$
 - There are more complex, necessary analysis
 - If task periods are harmonic (every period is an integer divisor of any larger period), then $U_b = 1$. This happens often in practice

EDF x RM

- RM is easier to implement in systems with limited number of priority levels
- RM is more transparent – easier to understand what is going on if something goes wrong (ex: overload)
 - I.e. if a task executes for longer than its prescribed worst-case time, higher priority tasks will be left untouched.

Different deadline assignments


- What happens if $D_i \neq T_i$?
- EDF still optimal
- Instead of RM, use DM – deadline monotonic (best among fixed priority algorithms)
- If $D_i < T_i$, utilization bound still works by changing periods to deadlines in the formula – however this is pessimistic

Different deadline assignments

- There exist exact schedulability analyses for both EDF and fixed priority that can be applied whatever the deadline – but they are pseudo-polynomial
- There are also analyses for asynchronous task sets – but exact analysis is exponential

Deadline Monotonic (DM)

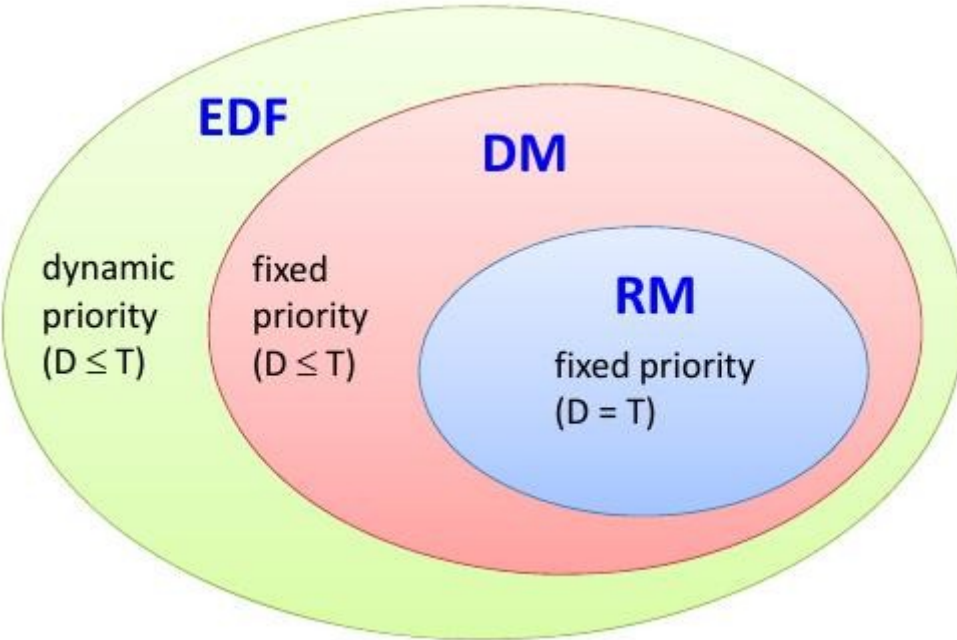
- Assign priorities according to the deadline
 - shorter deadline = higher priority
- Optimal when $D_i < T_i$ for fixed priority scheduling



- Problem with the utilization bound

$$U_p = \sum_{i=1}^n \frac{C_i}{D_i} = \frac{2}{3} + \frac{3}{6} = 1.16 > 1 \quad \text{but the task set is schedulable}$$

Optimality



Response Time Analysis (RTA)

- Proposed by Audsley 90
- Iterative solution for fixed priority systems
- Exact schedulability test
- Tasks ordered by decreasing priority
- $R_i^{(k)}$: worst-case response time for τ_i at step k

$$R_i^{(0)} = C_i + \sum_{j=1}^{i-1} C_j$$

$$R_i^{(k)} = C_i + \sum_{j=1}^{i-1} \left\lceil \frac{R_i^{(k-1)}}{T_j} \right\rceil C_j$$


Response Time Analysis (RTA)

- If the iteration converges \rightarrow fixed point represents the worst-case response time for the task
- key idea: $\left\lceil \frac{R_i^{(k-1)}}{T_j} \right\rceil$ is the max # of jobs of τ_j that interfere with τ_i

RTA: example

- τ_1 ($C_1 = 2, T_1 = 5$), τ_2 ($C_2 = 2, T_2 = 9$), τ_3 ($C_2 = 5, T_2 = 20$); $U = 0.872 > U_b(3) = 0.780$

$$R_3^{(0)} = C_3 + 1 \cdot C_1 + 1 \cdot C_2 = 9$$




RTA: example

- τ_1 ($C_1 = 2, T_1 = 5$), τ_2 ($C_2 = 2, T_1 = 9$), τ_3 ($C_2 = 5, T_1 = 20$); $U = 0.872 > U_b(3) = 0.780$

$$R_3^{(0)} = C_3 + 1 \cdot C_1 + 1 \cdot C_2 = 9$$

$$R_3^{(1)} = C_3 + 2 \cdot C_1 + 1 \cdot C_2 = 11$$




RTA: example

- τ_1 ($C_1 = 2, T_1 = 5$), τ_2 ($C_2 = 2, T_2 = 9$), τ_3 ($C_3 = 5, T_3 = 20$); $U = 0.872 > U_b(3) = 0.780$

$$R_3^{(0)} = C_3 + 1 \cdot C_1 + 1 \cdot C_2 = 9$$

$$R_3^{(1)} = C_3 + 2 \cdot C_1 + 1 \cdot C_2 = 11$$

$$R_3^{(2)} = C_3 + 3 \cdot C_1 + 2 \cdot C_2 = 15$$



RTA: example


- τ_1 ($C_1 = 2, T_1 = 5$), τ_2 ($C_2 = 2, T_2 = 9$), τ_3 ($C_3 = 5, T_3 = 20$); $U = 0.872 > U_b(3) = 0.780$

$$R_3^{(0)} = C_3 + 1 \cdot C_1 + 1 \cdot C_2 = 9$$

$$R_3^{(1)} = C_3 + 2 \cdot C_1 + 1 \cdot C_2 = 11$$

$$R_3^{(2)} = C_3 + 3 \cdot C_1 + 2 \cdot C_2 = 15$$

$$R_3^{(3)} = C_3 + 3 \cdot C_1 + 2 \cdot C_2 = 15 = R_3^{(2)}$$



Least Laxity First (LLF)

- Assign a higher priority to a task with smaller laxity
 - As EDF, it is optimal
- There is the need to know the current time to compute the priority

Review

- Periodic task model
- Uniprocessor RT scheduling
 - Task are independent
 - Cyclic executive
 - RM
 - EDF
 - RM
 - LLF

References

- Giorgio Buttazo. Real-time systems course. Scuola Superiore Santanna.
- Rodolfo Pellizzoni. Cyber-physical systems course. University of Waterloo.
- Farines, Silva, Oliveira. Sistemas de Tempo Real. DAS/UFSC. Julho 2000.
- G.C. Buttazzo: Hard Real-Time Computing Systems: Predictable Scheduling Algorithms and applications. Springer, 2004.